

task1

November 28, 2022

1

№1

:

```
[1]: !pip install -q tqdm
      !pip install --upgrade --no-cache-dir gdown
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Requirement already satisfied: gdown in `/usr/local/lib/python3.7/dist-packages` (4.4.0)

Collecting gdown

Downloading gdown-4.5.4-py3-none-any.whl (14 kB)

Requirement already satisfied: six in `/usr/local/lib/python3.7/dist-packages` (from gdown) (1.15.0)

Requirement already satisfied: tqdm in `/usr/local/lib/python3.7/dist-packages` (from gdown) (4.64.1)

Requirement already satisfied: beautifulsoup4 in `/usr/local/lib/python3.7/dist-packages` (from gdown) (4.6.3)

Requirement already satisfied: requests[socks] in `/usr/local/lib/python3.7/dist-packages` (from gdown) (2.23.0)

Requirement already satisfied: filelock in `/usr/local/lib/python3.7/dist-packages` (from gdown) (3.8.0)

Requirement already satisfied: chardet<4,>=3.0.2 in `/usr/local/lib/python3.7/dist-packages` (from requests[socks]->gdown) (3.0.4)

Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in `/usr/local/lib/python3.7/dist-packages` (from requests[socks]->gdown) (1.24.3)

Requirement already satisfied: idna<3,>=2.5 in `/usr/local/lib/python3.7/dist-packages` (from requests[socks]->gdown) (2.10)

Requirement already satisfied: certifi>=2017.4.17 in `/usr/local/lib/python3.7/dist-packages` (from requests[socks]->gdown) (2022.9.24)

Requirement already satisfied: PySocks!=1.5.7,>=1.5.6 in `/usr/local/lib/python3.7/dist-packages` (from requests[socks]->gdown) (1.7.1)

Installing collected packages: gdown

Attempting uninstall: gdown

Found existing installation: gdown 4.4.0

Uninstalling gdown-4.4.0:

Successfully uninstalled gdown-4.4.0
Successfully installed gdown-4.5.4

Google Drive :

```
[2]: from google.colab import drive
drive.mount('/content/drive', force_remount=True)
```

Mounted at /content/drive

, , (gdrive) :

```
[3]: EVALUATE_ONLY = True
TEST_ON_LARGE_DATASET = True
TISSUE_CLASSES = ('ADI', 'BACK', 'DEB', 'LYM', 'MUC', 'MUS', 'NORM', 'STR', 'TUM')
DATASETS_LINKS = {
    'train': '1XtQzVQ5XbrfxpLHJuLOXBGJ5U7CS-cLi',
    'train_small': '1qd45xXfDwdZjktLFwQb-et-mAaFeCzOR',
    'train_tiny': '1I-2Z0uXLd4QwhZQqltp817Kn3JOXgbui',
    'test': '1RfPou3pFKpuHDJZ-D9XDFzgvpUBF1Dr',
    'test_small': '1wbRsogOn7uG1HIPGLhyN-PMET2kdQ2lI',
    'test_tiny': '1viiB0s041CNsAK4itvX8PnYthJ-MDnQc'
}
```

:

```
[4]: from pathlib import Path
import numpy as np
from typing import List
from tqdm.notebook import tqdm
from time import sleep
from PIL import Image
import IPython.display
from sklearn.metrics import balanced_accuracy_score
import gdown
```

```
[5]: # extra imports

import torch
from torch import nn
from torch.utils.tensorboard import SummaryWriter
from torchvision.transforms.functional import to_pil_image
import os
from collections import namedtuple
import matplotlib.pyplot as plt
```

```
[6]: %load_ext tensorboard
```

1.0.1 Dataset

().

```
[7]: class Dataset:

    def __init__(self, name):
        self.name = name
        self.is_loaded = False
        url = f"https://drive.google.com/uc?
→export=download&confirm=pbef&id={DATASETS_LINKS[name]}"
        output = f'{name}.npz'
        gdown.download(url, output, quiet=False)
        print(f'Loading dataset {self.name} from npz.')
        print('before')
        np_obj = np.load(f'{name}.npz')
        print('after')
        self.images = np_obj['data']
        self.labels = np_obj['labels']
        self.n_files = self.images.shape[0]
        self.is_loaded = True
        print(f'Done. Dataset {name} consists of {self.n_files} images.')

    def image(self, i):
        # read i-th image in dataset and return it as numpy array
        if self.is_loaded:
            return self.images[i, :, :, :]

    def images_seq(self, n=None):
        # sequential access to images inside dataset (is needed for testing)
        for i in range(self.n_files if not n else n):
            yield self.image(i)

    def random_image_with_label(self):
        # get random image with label from dataset
        i = np.random.randint(self.n_files)
        return self.image(i), self.labels[i]

    def random_batch_with_labels(self, n):
        # create random batch of images with labels (is needed for training)
        indices = np.random.choice(self.n_files, n)
        imgs = []
        for i in indices:
            img = self.image(i)
            imgs.append(self.image(i))
```

```

        logits = np.array([self.labels[i] for i in indices])
        return np.stack(imgs), logits

    def image_with_label(self, i: int):
        # return i-th image with label from dataset
        return self.image(i), self.labels[i]

```

1.0.2 Dataset

```

[8]: d_train_tiny = Dataset('train_tiny')

img, lbl = d_train_tiny.random_image_with_label()
print()
print(f'Got numpy array of shape {img.shape}, and label with code {lbl}.')
print(f'Label code corresponds to {TISSUE_CLASSES[lbl]} class.')

pil_img = Image.fromarray(img)
IPython.display.display(pil_img)

```

Downloading...

From: <https://drive.google.com/uc?export=download&confirm=pbef&id=1I-2Z0uXLd4QwhZQQltp817Kn3J0Xgbui>

To: /content/train_tiny.npz

100%| | 105M/105M [00:00<00:00, 275MB/s]

Loading dataset train_tiny from npz.


before

after

Done. Dataset train_tiny consists of 900 images.

Got numpy array of shape (224, 224, 3), and label with code 1.

Label code corresponds to BACK class.



1.0.3 dataloader Pytorch

```
[9]: class TissueDataset(torch.utils.data.Dataset):
    def __init__(self, dataset: Dataset, mode: str, transforms=None):
        # mode : train, validation, full

        images, labels = dataset.images, dataset.labels
        self.transforms = transforms
        # LBL1
        if mode == 'full':
            self.samples = list(zip(images, labels))
        else:
            train_size = int(0.8 * len(images))
            val_size = len(images) - train_size
            train_dataset, val_dataset = torch.utils.data.
            random_split(list(zip(images, labels)), [train_size, val_size])
            if mode == 'train':
                self.samples = train_dataset
            else:
                self.samples = val_dataset

    def __getitem__(self, idx: int):
        image, label = self.samples[idx]
        if self.transforms:
            image = self.transforms(image)
```

```

        image = torch.tensor(image.transpose(2, 0, 1), dtype=torch.float32) / 255
        return image, label

    def __len__(self):
        return len(self.samples)

```


```

[10]: # example

data = TissueDataset(d_train_tiny, 'validation')
for img, label in data:
    plt.imshow(img.permute(1, 2, 0))
    print(img.size(), label)
    break

```

torch.Size([3, 224, 224]) 8



1.0.4 Metrics

, : 1. , 2. .

```

[11]: class Metrics:

```

```

    @staticmethod
    def accuracy(gt: List[int], pred: List[int]):
        assert len(gt) == len(pred), 'gt and prediction should be of equal_
→length'
        return sum(int(i[0] == i[1]) for i in zip(gt, pred)) / len(gt)

    @staticmethod
    def accuracy_balanced(gt: List[int], pred: List[int]):
        return balanced_accuracy_score(gt, pred)

    @staticmethod
    def print_all(gt: List[int], pred: List[int], info: str):
        print(f'metrics for {info}:')
        print('\t accuracy {:.4f}:'.format(Metrics.accuracy(gt, pred)))
        print('\t balanced accuracy {:.4f}:'.format(Metrics.
→accuracy_balanced(gt, pred)))

```

1.0.5 Model

```

,
.
save, load
.
, ,
, ,
,
,
, : 1. ; 2. ; 3.
; 4. - ( ); 6. , /
5. ( , ); 7.
( ( ); 8.
9. ; 10.
( ) 11. ..
.
.

```

```

[12]: from collections import Counter

class ConvBlock(nn.Module):
    def __init__(self, in_channels: int, out_channels: int):
        super(ConvBlock, self).__init__()

        self.convblock = nn.Sequential(

```

```

        nn.Conv2d(in_channels=in_channels, out_channels=out_channels,
        ↪kernel_size=3, stride=1, padding=1, bias=True),
        nn.BatchNorm2d(num_features=out_channels),
        nn.ReLU(inplace=True)
    )

    def forward(self, x):
        out = self.convblock(x)
        return out

class ConvNet(nn.Module):
    cfg = [16, "M", 32, "M", 64, "M", 128, "M", 256, "M"]
    def __init__(self, in_channels: int, num_classes: int):
        super().__init__()

        self.features = torch.nn.Sequential()
        last_output = in_channels
        for i in ConvNet.cfg:
            if isinstance(i, int):
                self.features.append(ConvBlock(last_output, i))
                last_output = i
            else:
                self.features.append(torch.nn.MaxPool2d(kernel_size=2))

        spatial_size = 224 / 2**(Counter(ConvNet.cfg)['M'])
        assert spatial_size == int(spatial_size) # has to be int
        spatial_size = int(spatial_size)
        self.classification_head = nn.Sequential(
            nn.Flatten(),
            nn.Linear(last_output * spatial_size * spatial_size, 128),
            nn.BatchNorm1d(128),
            nn.ReLU(),
            nn.Dropout(p=0.5),
            nn.Linear(128, num_classes)
        )

    def forward(self, x):
        out = self.features(x)
        return self.classification_head(out)

```

```

[13]: device = torch.device('cpu')
if torch.cuda.is_available():
    device = torch.device('cuda', 0)

net = ConvNet(3, 9)
net.to(device)

```



```

[13]: ConvNet(
  (features): Sequential(
    (0): ConvBlock(
      (convblock): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
      )
    )
    (1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (2): ConvBlock(
      (convblock): Sequential(
        (0): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
      )
    )
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (4): ConvBlock(
      (convblock): Sequential(
        (0): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
      )
    )
    (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (6): ConvBlock(
      (convblock): Sequential(
        (0): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
      )
    )
    (7): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (8): ConvBlock(
      (convblock): Sequential(
        (0): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        (2): ReLU(inplace=True)
    )
)
(9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
)
(classification_head): Sequential(
  (0): Flatten(start_dim=1, end_dim=-1)
  (1): Linear(in_features=12544, out_features=128, bias=True)
  (2): BatchNorm1d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (3): ReLU()
  (4): Dropout(p=0.5, inplace=False)
  (5): Linear(in_features=128, out_features=9, bias=True)
)
)

```

```

[21]: class Model:
    def __init__(self, cfg):
        self.cfg = cfg
        self.device = cfg.device
        self.out_dir = os.path.join(cfg.out_dir, cfg.model_name)
        os.makedirs(self.out_dir, exist_ok=True)
        log_dir = os.path.join(self.out_dir, "runs")
        os.makedirs(log_dir, exist_ok=True)

        self.writer = SummaryWriter(log_dir=log_dir)

        #LBL5
        if cfg.num_pretrained_epoch > 0:
            self.model = ConvNet(3, 9)
            weights_dir = os.path.join(self.out_dir, "weights")
            self.load_pretrained(os.path.join(weights_dir, f"{self.cfg.
→model_name}_epoch_{cfg.num_pretrained_epoch}.pth"))
        else:
            self.model = ConvNet(3, 9)

        self.criterion = nn.CrossEntropyLoss()

    def save(self, path: str):
        torch.save(self.model.state_dict(), path)

    def load(self, name: str):
        # https://drive.google.com/file/d/1nrASFxOMWfRB7f9FLc322uM-I-l8UHL6/view?
→usp=sharing

```

```

name_to_id_dict = '1nrASFx0MwfRB7f9FLc322uM-I-18UHL6'
output = f'{name}.pth'
gdown.download(f'https://drive.google.com/uc?id={name_to_id_dict}',
↳output, quiet=False)
self.model.load_state_dict(torch.load(output, map_location='cpu'))

def load_pretrained(self, path: str):
    self.model.load_state_dict(torch.load(path))

def train(self, train_ds: Dataset):
    self.model.to(self.device)
    params = [p for p in self.model.parameters() if p.requires_grad]
    optimizer = torch.optim.Adam(params, self.cfg.lr)
    lr_scheduler = torch.optim.lr_scheduler.MultiStepLR(optimizer,
↳milestones=self.cfg.milestones)

    train_dl, val_dl = self.get_train_dataloaders(
        train_ds,
        self.cfg.batch_size, self.cfg.batch_size_val, self.cfg.num_workers
    )

    weights_dir = os.path.join(self.out_dir, "weights")
    os.makedirs(weights_dir, exist_ok=True)

    for epoch in range(1, self.cfg.epochs + 1):
        print(f"Epoch: {epoch}")

        logs = self.train_epoch(optimizer, train_dl, epoch)
        for key, value in logs.items():
            self.writer.add_scalar(key, value, epoch)

        if lr_scheduler is not None:
            lr_scheduler.step()

        #LBL3
        metrics = self.evaluate(val_dl)
        for key, value in metrics.items():
            self.writer.add_scalar(f"val/{key}", value, epoch)

        if epoch % 4 == 0:
            self.save(os.path.join(weights_dir, f"{self.cfg.
↳model_name}_epoch_{epoch}.pth"))

        #LBL2
        self.save(os.path.join(weights_dir, f"{self.cfg.model_name}.pth"))

```

```

def train_epoch(self, optimizer, train_dl, epoch):
    self.model.train()

    lr_scheduler = None
    if epoch == 1:
        warmup_factor = 1e-3
        warmup_iters = len(train_dl)
        #print(warmup_iters, 'warmup_iters')

        lr_scheduler = torch.optim.lr_scheduler.LinearLR(
            optimizer, start_factor=warmup_factor, total_iters=warmup_iters,
            verbose=True
        )

    for images, targets in tqdm(train_dl):
        optimizer.zero_grad()
        loss = self.criterion(self.model(images.to(self.device)), targets.
→to(self.device))

        loss.backward()
        optimizer.step()

        if lr_scheduler is not None:
            lr_scheduler.step()

    #LBL4
    print(f"train loss: {loss}")
    return {"loss": loss.detach()}

def evaluate(self, val_dl):
    self.model.eval()

    prediction, target = [], []
    for images, targets in tqdm(val_dl):
        outputs = self.model(images.to(self.device)).to('cpu')
        target += list(targets.numpy())
        _, predicted = torch.max(outputs, 1)
        prediction += list(predicted.detach().numpy())

    acc = Metrics.accuracy(target, prediction)
    balanced_acc = Metrics.accuracy_balanced(target, prediction)

    #LBL4
    print(f"val accuracy: {acc}")

```

```

        return {"accuracy": acc, "balanced_acc": balanced_acc}

    def get_train_dataloaders(self, ds_numpy, batch_size, batch_size_val,
    ↪num_workers):
        train_ds = TissueDataset(ds_numpy, "train")
        val_ds = TissueDataset(ds_numpy, "validation")

        train_dl = torch.utils.data.DataLoader(
            train_ds,
            batch_size=batch_size,
            num_workers=num_workers,
            shuffle=True
        )
        val_dl = torch.utils.data.DataLoader(
            val_ds,
            batch_size=batch_size_val,
            num_workers=num_workers,
            shuffle=False,
        )
        return train_dl, val_dl

    def test_on_dataset(self, dataset: Dataset, limit=None):
        self.model.eval()
        self.model.to('cpu')
        predictions = []
        n = dataset.n_files if not limit else int(dataset.n_files * limit)
        for img in tqdm(dataset.images_seq(n), total=n):
            predictions.append(self.test_on_image(img))
        return predictions

    def test_on_image(self, img: np.ndarray):
        img = torch.tensor(img.transpose(2, 0, 1), dtype=torch.float32) / 255
        _, prediction = torch.max(self.model(img.unsqueeze(0)), 1)
        return prediction.numpy()

```

1.0.6

'train_small' 'test_small'.

```

[17]: d_train = Dataset('train')
      d_test = Dataset('test')

```

```

Downloading...
From: https://drive.google.com/uc?export=download&confirm=pbef&id=1XtQzVQ5Xbrfxp
LHJuLOXBGJ5U7CS-cLi
To: /content/train.npz
100%|          | 2.10G/2.10G [00:54<00:00, 38.2MB/s]

Loading dataset train from npz.
before
after
Done. Dataset train consists of 18000 images.

Downloading...
From: https://drive.google.com/uc?export=download&confirm=pbef&id=1RfPou3pFKpuHD
JZ-D9XDFzgwpUBFlDr
To: /content/test.npz
100%|          | 525M/525M [00:12<00:00, 42.0MB/s]

Loading dataset test from npz.
before
after
Done. Dataset test consists of 4500 images.

```

```

[18]: cfg_dict = {
    "out_dir": Path('drive/MyDrive/nn_msu/'),
    "batch_size": 64,
    "batch_size_val": 64,
    "num_workers": 2,
    "model_name": 'ConvNet',
    "num_pretrained_epoch": 0,
    "device": device,
    "epochs": 10,
    "lr": 0.001,
    "weight_decay": 0.0001,
    "milestones": [6, 8, 9],
}

cfg = namedtuple("Config", cfg_dict.keys())(**cfg_dict)

```

```

[19]: EVALUATE_ONLY = False

```

1.0.7 Train

```

[22]: model = Model(cfg)
    if not EVALUATE_ONLY:
        model.train(d_train)
        model.save('ConvNetModel')
    else:
        #todo: your link goes here

```

```
model.load('ConvNetModel')
```

Epoch: 1

Adjusting learning rate of group 0 to 1.0000e-06.

0%| | 0/225 [00:00<?, ?it/s]

Adjusting learning rate of group 0 to 5.4400e-06.
Adjusting learning rate of group 0 to 9.8800e-06.
Adjusting learning rate of group 0 to 1.4320e-05.
Adjusting learning rate of group 0 to 1.8760e-05.
Adjusting learning rate of group 0 to 2.3200e-05.
Adjusting learning rate of group 0 to 2.7640e-05.
Adjusting learning rate of group 0 to 3.2080e-05.
Adjusting learning rate of group 0 to 3.6520e-05.
Adjusting learning rate of group 0 to 4.0960e-05.
Adjusting learning rate of group 0 to 4.5400e-05.
Adjusting learning rate of group 0 to 4.9840e-05.
Adjusting learning rate of group 0 to 5.4280e-05.
Adjusting learning rate of group 0 to 5.8720e-05.
Adjusting learning rate of group 0 to 6.3160e-05.
Adjusting learning rate of group 0 to 6.7600e-05.
Adjusting learning rate of group 0 to 7.2040e-05.
Adjusting learning rate of group 0 to 7.6480e-05.
Adjusting learning rate of group 0 to 8.0920e-05.
Adjusting learning rate of group 0 to 8.5360e-05.
Adjusting learning rate of group 0 to 8.9800e-05.
Adjusting learning rate of group 0 to 9.4240e-05.
Adjusting learning rate of group 0 to 9.8680e-05.
Adjusting learning rate of group 0 to 1.0312e-04.
Adjusting learning rate of group 0 to 1.0756e-04.
Adjusting learning rate of group 0 to 1.1200e-04.
Adjusting learning rate of group 0 to 1.1644e-04.
Adjusting learning rate of group 0 to 1.2088e-04.
Adjusting learning rate of group 0 to 1.2532e-04.
Adjusting learning rate of group 0 to 1.2976e-04.
Adjusting learning rate of group 0 to 1.3420e-04.
Adjusting learning rate of group 0 to 1.3864e-04.
Adjusting learning rate of group 0 to 1.4308e-04.
Adjusting learning rate of group 0 to 1.4752e-04.
Adjusting learning rate of group 0 to 1.5196e-04.
Adjusting learning rate of group 0 to 1.5640e-04.
Adjusting learning rate of group 0 to 1.6084e-04.
Adjusting learning rate of group 0 to 1.6528e-04.
Adjusting learning rate of group 0 to 1.6972e-04.
Adjusting learning rate of group 0 to 1.7416e-04.
Adjusting learning rate of group 0 to 1.7860e-04.

Adjusting learning rate of group 0 to 1.8304e-04.
Adjusting learning rate of group 0 to 1.8748e-04.
Adjusting learning rate of group 0 to 1.9192e-04.
Adjusting learning rate of group 0 to 1.9636e-04.
Adjusting learning rate of group 0 to 2.0080e-04.
Adjusting learning rate of group 0 to 2.0524e-04.
Adjusting learning rate of group 0 to 2.0968e-04.
Adjusting learning rate of group 0 to 2.1412e-04.
Adjusting learning rate of group 0 to 2.1856e-04.
Adjusting learning rate of group 0 to 2.2300e-04.
Adjusting learning rate of group 0 to 2.2744e-04.
Adjusting learning rate of group 0 to 2.3188e-04.
Adjusting learning rate of group 0 to 2.3632e-04.
Adjusting learning rate of group 0 to 2.4076e-04.
Adjusting learning rate of group 0 to 2.4520e-04.
Adjusting learning rate of group 0 to 2.4964e-04.
Adjusting learning rate of group 0 to 2.5408e-04.
Adjusting learning rate of group 0 to 2.5852e-04.
Adjusting learning rate of group 0 to 2.6296e-04.
Adjusting learning rate of group 0 to 2.6740e-04.
Adjusting learning rate of group 0 to 2.7184e-04.
Adjusting learning rate of group 0 to 2.7628e-04.
Adjusting learning rate of group 0 to 2.8072e-04.
Adjusting learning rate of group 0 to 2.8516e-04.
Adjusting learning rate of group 0 to 2.8960e-04.
Adjusting learning rate of group 0 to 2.9404e-04.
Adjusting learning rate of group 0 to 2.9848e-04.
Adjusting learning rate of group 0 to 3.0292e-04.
Adjusting learning rate of group 0 to 3.0736e-04.
Adjusting learning rate of group 0 to 3.1180e-04.
Adjusting learning rate of group 0 to 3.1624e-04.
Adjusting learning rate of group 0 to 3.2068e-04.
Adjusting learning rate of group 0 to 3.2512e-04.
Adjusting learning rate of group 0 to 3.2956e-04.
Adjusting learning rate of group 0 to 3.3400e-04.
Adjusting learning rate of group 0 to 3.3844e-04.
Adjusting learning rate of group 0 to 3.4288e-04.
Adjusting learning rate of group 0 to 3.4732e-04.
Adjusting learning rate of group 0 to 3.5176e-04.
Adjusting learning rate of group 0 to 3.5620e-04.
Adjusting learning rate of group 0 to 3.6064e-04.
Adjusting learning rate of group 0 to 3.6508e-04.
Adjusting learning rate of group 0 to 3.6952e-04.
Adjusting learning rate of group 0 to 3.7396e-04.
Adjusting learning rate of group 0 to 3.7840e-04.
Adjusting learning rate of group 0 to 3.8284e-04.
Adjusting learning rate of group 0 to 3.8728e-04.
Adjusting learning rate of group 0 to 3.9172e-04.

Adjusting learning rate of group 0 to 3.9616e-04.
Adjusting learning rate of group 0 to 4.0060e-04.
Adjusting learning rate of group 0 to 4.0504e-04.
Adjusting learning rate of group 0 to 4.0948e-04.
Adjusting learning rate of group 0 to 4.1392e-04.
Adjusting learning rate of group 0 to 4.1836e-04.
Adjusting learning rate of group 0 to 4.2280e-04.
Adjusting learning rate of group 0 to 4.2724e-04.
Adjusting learning rate of group 0 to 4.3168e-04.
Adjusting learning rate of group 0 to 4.3612e-04.
Adjusting learning rate of group 0 to 4.4056e-04.
Adjusting learning rate of group 0 to 4.4500e-04.
Adjusting learning rate of group 0 to 4.4944e-04.
Adjusting learning rate of group 0 to 4.5388e-04.
Adjusting learning rate of group 0 to 4.5832e-04.
Adjusting learning rate of group 0 to 4.6276e-04.
Adjusting learning rate of group 0 to 4.6720e-04.
Adjusting learning rate of group 0 to 4.7164e-04.
Adjusting learning rate of group 0 to 4.7608e-04.
Adjusting learning rate of group 0 to 4.8052e-04.
Adjusting learning rate of group 0 to 4.8496e-04.
Adjusting learning rate of group 0 to 4.8940e-04.
Adjusting learning rate of group 0 to 4.9384e-04.
Adjusting learning rate of group 0 to 4.9828e-04.
Adjusting learning rate of group 0 to 5.0272e-04.
Adjusting learning rate of group 0 to 5.0716e-04.
Adjusting learning rate of group 0 to 5.1160e-04.
Adjusting learning rate of group 0 to 5.1604e-04.
Adjusting learning rate of group 0 to 5.2048e-04.
Adjusting learning rate of group 0 to 5.2492e-04.
Adjusting learning rate of group 0 to 5.2936e-04.
Adjusting learning rate of group 0 to 5.3380e-04.
Adjusting learning rate of group 0 to 5.3824e-04.
Adjusting learning rate of group 0 to 5.4268e-04.
Adjusting learning rate of group 0 to 5.4712e-04.
Adjusting learning rate of group 0 to 5.5156e-04.
Adjusting learning rate of group 0 to 5.5600e-04.
Adjusting learning rate of group 0 to 5.6044e-04.
Adjusting learning rate of group 0 to 5.6488e-04.
Adjusting learning rate of group 0 to 5.6932e-04.
Adjusting learning rate of group 0 to 5.7376e-04.
Adjusting learning rate of group 0 to 5.7820e-04.
Adjusting learning rate of group 0 to 5.8264e-04.
Adjusting learning rate of group 0 to 5.8708e-04.
Adjusting learning rate of group 0 to 5.9152e-04.
Adjusting learning rate of group 0 to 5.9596e-04.
Adjusting learning rate of group 0 to 6.0040e-04.
Adjusting learning rate of group 0 to 6.0484e-04.

Adjusting learning rate of group 0 to 6.0928e-04.
Adjusting learning rate of group 0 to 6.1372e-04.
Adjusting learning rate of group 0 to 6.1816e-04.
Adjusting learning rate of group 0 to 6.2260e-04.
Adjusting learning rate of group 0 to 6.2704e-04.
Adjusting learning rate of group 0 to 6.3148e-04.
Adjusting learning rate of group 0 to 6.3592e-04.
Adjusting learning rate of group 0 to 6.4036e-04.
Adjusting learning rate of group 0 to 6.4480e-04.
Adjusting learning rate of group 0 to 6.4924e-04.
Adjusting learning rate of group 0 to 6.5368e-04.
Adjusting learning rate of group 0 to 6.5812e-04.
Adjusting learning rate of group 0 to 6.6256e-04.
Adjusting learning rate of group 0 to 6.6700e-04.
Adjusting learning rate of group 0 to 6.7144e-04.
Adjusting learning rate of group 0 to 6.7588e-04.
Adjusting learning rate of group 0 to 6.8032e-04.
Adjusting learning rate of group 0 to 6.8476e-04.
Adjusting learning rate of group 0 to 6.8920e-04.
Adjusting learning rate of group 0 to 6.9364e-04.
Adjusting learning rate of group 0 to 6.9808e-04.
Adjusting learning rate of group 0 to 7.0252e-04.
Adjusting learning rate of group 0 to 7.0696e-04.
Adjusting learning rate of group 0 to 7.1140e-04.
Adjusting learning rate of group 0 to 7.1584e-04.
Adjusting learning rate of group 0 to 7.2028e-04.
Adjusting learning rate of group 0 to 7.2472e-04.
Adjusting learning rate of group 0 to 7.2916e-04.
Adjusting learning rate of group 0 to 7.3360e-04.
Adjusting learning rate of group 0 to 7.3804e-04.
Adjusting learning rate of group 0 to 7.4248e-04.
Adjusting learning rate of group 0 to 7.4692e-04.
Adjusting learning rate of group 0 to 7.5136e-04.
Adjusting learning rate of group 0 to 7.5580e-04.
Adjusting learning rate of group 0 to 7.6024e-04.
Adjusting learning rate of group 0 to 7.6468e-04.
Adjusting learning rate of group 0 to 7.6912e-04.
Adjusting learning rate of group 0 to 7.7356e-04.
Adjusting learning rate of group 0 to 7.7800e-04.
Adjusting learning rate of group 0 to 7.8244e-04.
Adjusting learning rate of group 0 to 7.8688e-04.
Adjusting learning rate of group 0 to 7.9132e-04.
Adjusting learning rate of group 0 to 7.9576e-04.
Adjusting learning rate of group 0 to 8.0020e-04.
Adjusting learning rate of group 0 to 8.0464e-04.
Adjusting learning rate of group 0 to 8.0908e-04.
Adjusting learning rate of group 0 to 8.1352e-04.
Adjusting learning rate of group 0 to 8.1796e-04.

Adjusting learning rate of group 0 to 8.2240e-04.
Adjusting learning rate of group 0 to 8.2684e-04.
Adjusting learning rate of group 0 to 8.3128e-04.
Adjusting learning rate of group 0 to 8.3572e-04.
Adjusting learning rate of group 0 to 8.4016e-04.
Adjusting learning rate of group 0 to 8.4460e-04.
Adjusting learning rate of group 0 to 8.4904e-04.
Adjusting learning rate of group 0 to 8.5348e-04.
Adjusting learning rate of group 0 to 8.5792e-04.
Adjusting learning rate of group 0 to 8.6236e-04.
Adjusting learning rate of group 0 to 8.6680e-04.
Adjusting learning rate of group 0 to 8.7124e-04.
Adjusting learning rate of group 0 to 8.7568e-04.
Adjusting learning rate of group 0 to 8.8012e-04.
Adjusting learning rate of group 0 to 8.8456e-04.
Adjusting learning rate of group 0 to 8.8900e-04.
Adjusting learning rate of group 0 to 8.9344e-04.
Adjusting learning rate of group 0 to 8.9788e-04.
Adjusting learning rate of group 0 to 9.0232e-04.
Adjusting learning rate of group 0 to 9.0676e-04.
Adjusting learning rate of group 0 to 9.1120e-04.
Adjusting learning rate of group 0 to 9.1564e-04.
Adjusting learning rate of group 0 to 9.2008e-04.
Adjusting learning rate of group 0 to 9.2452e-04.
Adjusting learning rate of group 0 to 9.2896e-04.
Adjusting learning rate of group 0 to 9.3340e-04.
Adjusting learning rate of group 0 to 9.3784e-04.
Adjusting learning rate of group 0 to 9.4228e-04.
Adjusting learning rate of group 0 to 9.4672e-04.
Adjusting learning rate of group 0 to 9.5116e-04.
Adjusting learning rate of group 0 to 9.5560e-04.
Adjusting learning rate of group 0 to 9.6004e-04.
Adjusting learning rate of group 0 to 9.6448e-04.
Adjusting learning rate of group 0 to 9.6892e-04.
Adjusting learning rate of group 0 to 9.7336e-04.
Adjusting learning rate of group 0 to 9.7780e-04.
Adjusting learning rate of group 0 to 9.8224e-04.
Adjusting learning rate of group 0 to 9.8668e-04.
Adjusting learning rate of group 0 to 9.9112e-04.
Adjusting learning rate of group 0 to 9.9556e-04.
Adjusting learning rate of group 0 to 1.0000e-03.
train loss: 0.5640280842781067

0%| | 0/57 [00:00<?, ?it/s]

val accuracy: 0.7813888888888889
Epoch: 2

```
0%|          | 0/225 [00:00<?, ?it/s]

train loss: 0.3132307529449463

0%|          | 0/57 [00:00<?, ?it/s]

val accuracy: 0.5336111111111111
Epoch: 3

0%|          | 0/225 [00:00<?, ?it/s]

train loss: 0.274355947971344

0%|          | 0/57 [00:00<?, ?it/s]

val accuracy: 0.4538888888888889
Epoch: 4

0%|          | 0/225 [00:00<?, ?it/s]

train loss: 0.35548290610313416

0%|          | 0/57 [00:00<?, ?it/s]

val accuracy: 0.8763888888888889
Epoch: 5

0%|          | 0/225 [00:00<?, ?it/s]

train loss: 0.2006574273109436

0%|          | 0/57 [00:00<?, ?it/s]

val accuracy: 0.8108333333333333
Epoch: 6

0%|          | 0/225 [00:00<?, ?it/s]

train loss: 0.15478897094726562

0%|          | 0/57 [00:00<?, ?it/s]

val accuracy: 0.8591666666666666
Epoch: 7

0%|          | 0/225 [00:00<?, ?it/s]
```

```
train loss: 0.06800190359354019
0%|          | 0/57 [00:00<?, ?it/s]
```

```
val accuracy: 0.9869444444444444
Epoch: 8
0%|          | 0/225 [00:00<?, ?it/s]
```

```
train loss: 0.02574928291141987
0%|          | 0/57 [00:00<?, ?it/s]
```

```
val accuracy: 0.9911111111111112
Epoch: 9
0%|          | 0/225 [00:00<?, ?it/s]
```

```
train loss: 0.028608273714780807
0%|          | 0/57 [00:00<?, ?it/s]
```

```
val accuracy: 0.99
Epoch: 10
0%|          | 0/225 [00:00<?, ?it/s]
```

```
train loss: 0.07655631750822067
0%|          | 0/57 [00:00<?, ?it/s]
```

```
val accuracy: 0.9905555555555555
```

1.0.8 Test

```
[23]: EVALUATE_ONLY = True

model = Model(cfg)
if not EVALUATE_ONLY:
    model.train(d_train)
    model.save('ConvNetModel')
else:
    #todo: your link goes here
    model.load('ConvNetModel')
```

Downloading...

From: <https://drive.google.com/uc?id=1nrASFx0MWfRB7f9FLc322uM-I-18UHL6>

```
To: /content/ConvNetModel.pth
100%|          | 8.02M/8.02M [00:00<00:00, 177MB/s]
```

:

```
[24]: # evaluating model on 10% of test dataset
pred_1 = model.test_on_dataset(d_test, limit=0.1)
Metrics.print_all(d_test.labels[:len(pred_1)], pred_1, '10% of test')
```

```
0%|          | 0/450 [00:00<?, ?it/s]
```

metrics for 10% of test:

accuracy 0.9956:

balanced accuracy 0.9956:

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1987:

UserWarning: y_pred contains classes not in y_true

warnings.warn("y_pred contains classes not in y_true")

:

```
[25]: # evaluating model on full test dataset (may take time)
if TEST_ON_LARGE_DATASET:
    pred_2 = model.test_on_dataset(d_test)
    Metrics.print_all(d_test.labels, pred_2, 'test')
```

```
0%|          | 0/4500 [00:00<?, ?it/s]
```

metrics for test:

accuracy 0.9729:

balanced accuracy 0.9729:

```
[26]: #LBL6
%tensorboard --logdir drive/MyDrive/nn_msu/ConvNet/runs
```

<IPython.core.display.Javascript object>

.

, ..

,

..

pdf (->)

pdf .

1.0.9

test_tiny,

(2%) test.

```
[28]: final_model = Model(cfg)
final_model.load('best')
d_test_tiny = Dataset('test_tiny')
pred = model.test_on_dataset(d_test_tiny)
Metrics.print_all(d_test_tiny.labels, pred, 'test-tiny')
```

Downloading...

From: <https://drive.google.com/uc?id=1nrASFx0MWfRB7f9FLc322uM-I-18UHL6>

To: /content/best.pth

100%| | 8.02M/8.02M [00:00<00:00, 62.2MB/s]

Downloading...

From: <https://drive.google.com/uc?export=download&confirm=pbef&id=1viiB0s041CNsAK4itvX8PnYthJ-MDnQc>

To: /content/test_tiny.npz

100%| | 10.6M/10.6M [00:00<00:00, 134MB/s]

Loading dataset test_tiny from npz.

before

after

Done. Dataset test_tiny consists of 90 images.

0%| | 0/90 [00:00<?, ?it/s]

metrics for test-tiny:

accuracy 0.9333:

balanced accuracy 0.9333:

Google Drive.

```
[29]: drive.flush_and_unmount()
```

2

”

”

2.0.1

timeit

:

```
[ ]: import timeit

def factorial(n):
    res = 1
```

```

    for i in range(1, n + 1):
        res *= i
    return res

def f():
    return factorial(n=1000)

n_runs = 128
print(f'Function f is caluclated {n_runs} times in {timeit.timeit(f,
↳number=n_runs)}s.')

```

2.0.2 Scikit-learn

learn.org/stable/). MNIST scikit-learn (<https://scikit-learn.org/stable/>). SVM:

```

[ ]: # Standard scientific Python imports
import matplotlib.pyplot as plt

# Import datasets, classifiers and performance metrics
from sklearn import datasets, svm, metrics
from sklearn.model_selection import train_test_split

# The digits dataset
digits = datasets.load_digits()

# The data that we are interested in is made of 8x8 images of digits, let's
# have a look at the first 4 images, stored in the `images` attribute of the
# dataset. If we were working from image files, we could load them using
# matplotlib.pyplot.imread. Note that each image must have the same size. For
↳these
# images, we know which digit they represent: it is given in the 'target' of
# the dataset.
_, axes = plt.subplots(2, 4)
images_and_labels = list(zip(digits.images, digits.target))
for ax, (image, label) in zip(axes[0, :], images_and_labels[:4]):
    ax.set_axis_off()
    ax.imshow(image, cmap=plt.cm.gray_r, interpolation='nearest')
    ax.set_title('Training: %i' % label)

# To apply a classifier on this data, we need to flatten the image, to
# turn the data in a (samples, feature) matrix:
n_samples = len(digits.images)
data = digits.images.reshape((n_samples, -1))

```



```

# Create a classifier: a support vector classifier
classifier = svm.SVC(gamma=0.001)

# Split data into train and test subsets
X_train, X_test, y_train, y_test = train_test_split(
    data, digits.target, test_size=0.5, shuffle=False)

# We learn the digits on the first half of the digits
classifier.fit(X_train, y_train)

# Now predict the value of the digit on the second half:
predicted = classifier.predict(X_test)

images_and_predictions = list(zip(digits.images[n_samples // 2:], predicted))
for ax, (image, prediction) in zip(axes[1, :], images_and_predictions[:4]):
    ax.set_axis_off()
    ax.imshow(image, cmap=plt.cm.gray_r, interpolation='nearest')
    ax.set_title('Prediction: %i' % prediction)

print("Classification report for classifier %s:\n%s\n"
      % (classifier, metrics.classification_report(y_test, predicted)))
disp = metrics.plot_confusion_matrix(classifier, X_test, y_test)
disp.figure_.suptitle("Confusion Matrix")
print("Confusion matrix:\n%s" % disp.confusion_matrix)

plt.show()

```

2.0.3 Scikit-image

, [scikit-image \(https://scikit-image.org/\)](https://scikit-image.org/), [numpy](#), [Canny edge detector](#).

```

[ ]: import numpy as np
import matplotlib.pyplot as plt
from scipy import ndimage as ndi

from skimage import feature

# Generate noisy image of a square
im = np.zeros((128, 128))
im[32:-32, 32:-32] = 1

im = ndi.rotate(im, 15, mode='constant')
im = ndi.gaussian_filter(im, 4)
im += 0.2 * np.random.random(im.shape)

```

```

# Compute the Canny filter for two values of sigma
edges1 = feature.canny(im)
edges2 = feature.canny(im, sigma=3)

# display results
fig, (ax1, ax2, ax3) = plt.subplots(nrows=1, ncols=3, figsize=(8, 3),
                                     sharex=True, sharey=True)

ax1.imshow(im, cmap=plt.cm.gray)
ax1.axis('off')
ax1.set_title('noisy image', fontsize=20)

ax2.imshow(edges1, cmap=plt.cm.gray)
ax2.axis('off')
ax2.set_title(r'Canny filter, $\sigma=1$', fontsize=20)

ax3.imshow(edges2, cmap=plt.cm.gray)
ax3.axis('off')
ax3.set_title(r'Canny filter, $\sigma=3$', fontsize=20)

fig.tight_layout()

plt.show()

```

2.0.4 Tensorflow 2

Tensorflow 2.

MNIST.

```

[ ]: # Install TensorFlow

import tensorflow as tf

mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

```

```

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)

model.evaluate(x_test, y_test, verbose=2)

```

GPU TPU. " " -> " " Google Colab

Tensorflow 2

<https://www.tensorflow.org/tutorials?hl=ru>.

Tensorflow 2.

TensorFlow 2. , (

), : <https://stanford.edu/~shervine/blog/keras-how-to-generate-data-on-the-fly>.

2.0.5 Numba

for python , JIT- Numba

(<https://numba.pydata.org/>). Numba Google Colab : 1.

https://colab.research.google.com/github/cbernet/maldives/blob/master/numba/numba_cuda.ipynb

2. https://colab.research.google.com/github/evaneschneider/parallel-programming/blob/master/COMPASS_gpu_intro.ipynb

Numba , , ,

Numba . Numba .

2.0.6 zip Google Drive

zip

zip . ,

Google Drive.

2 , tmp PROJECT_DIR, tmp tmp.zip.

```

[ ]: PROJECT_DIR = "/dev/prak_nn_1/"
arr1 = np.random.rand(100, 100, 3) * 255
arr2 = np.random.rand(100, 100, 3) * 255

img1 = Image.fromarray(arr1.astype('uint8'))
img2 = Image.fromarray(arr2.astype('uint8'))

p = "/content/drive/MyDrive/" + PROJECT_DIR

```

```

if not (Path(p) / 'tmp').exists():
    (Path(p) / 'tmp').mkdir()

img1.save(str(Path(p) / 'tmp' / 'img1.png'))
img2.save(str(Path(p) / 'tmp' / 'img2.png'))

%cd $p
!zip -r "tmp.zip" "tmp"

```

```

tmp.zip      tmp2 PROJECT_DIR.      tmp2      tmp,
2           .

```

```

[ ]: p = "/content/drive/MyDrive/" + PROJECT_DIR
%cd $p
!unzip -uq "tmp.zip" -d "tmp2"

```